

Artificial Intelligence, Machine Learning, and Retrieval-Augmented Generation

Introduction to AI (Section 1)

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Machine Learning (Section 1)

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Deep Learning (Section 1)

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Natural Language Processing (Section 1)

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

Large Language Models (Section 1)

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Embeddings (Section 1)

Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space.

Vector Databases (Section 1)

Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems.

Chunking (Section 1)

Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance.

Retrieval-Augmented Generation (Section 1)

RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context.

Hybrid Search (Section 1)

Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search

combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance.

Query Rewriting (Section 1)

Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems.

Reranking (Section 1)

Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking.

Conversation Memory (Section 1)

Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations.

Evaluation (Section 1)

RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction.

Future of RAG (Section 1)

Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities.

Introduction to AI (Section 2)

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Machine Learning (Section 2)

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Deep Learning (Section 2)

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Natural Language Processing (Section 2)

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

Large Language Models (Section 2)

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Embeddings (Section 2)

Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space.

Vector Databases (Section 2)

Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems.

Chunking (Section 2)

Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance.

Retrieval-Augmented Generation (Section 2)

RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context.

Hybrid Search (Section 2)

Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance.

Query Rewriting (Section 2)

Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems.

Reranking (Section 2)

Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking.

Conversation Memory (Section 2)

Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations.

Evaluation (Section 2)

RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction.

Future of RAG (Section 2)

Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities.

Introduction to AI (Section 3)

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Machine Learning (Section 3)

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Deep Learning (Section 3)

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Natural Language Processing (Section 3)

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

Large Language Models (Section 3)

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Embeddings (Section 3)

Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space.

Vector Databases (Section 3)

Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems.

Chunking (Section 3)

Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance.

Retrieval-Augmented Generation (Section 3)

RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context.

Hybrid Search (Section 3)

Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance.

Query Rewriting (Section 3)

Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems.

Reranking (Section 3)

Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking.

Conversation Memory (Section 3)

Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations.

Evaluation (Section 3)

RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction.

Future of RAG (Section 3)

Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities.

Introduction to AI (Section 4)

Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven. Artificial Intelligence (AI) is a branch of computer science focused on creating systems capable of performing tasks that normally require human intelligence. These tasks include reasoning, learning, planning, perception, and language understanding. AI systems can be rule-based or data-driven.

Machine Learning (Section 4)

Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning. Machine Learning is a subset of AI where systems learn patterns from data rather than being explicitly programmed. Common paradigms include supervised learning, unsupervised learning, and reinforcement learning.

Deep Learning (Section 4)

Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing. Deep Learning uses neural networks with multiple layers to learn complex representations of data. Applications include computer vision, speech recognition, and natural language processing.

Natural Language Processing (Section 4)

NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures. NLP enables computers to understand, interpret, and generate human language. Modern NLP systems are often powered by transformer-based architectures.

Large Language Models (Section 4)

Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast

amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation. Large Language Models are trained on vast amounts of text data and can perform tasks such as question answering, summarization, translation, and content generation.

Embeddings (Section 4)

Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space. Embeddings convert text into numerical vectors that capture semantic meaning. Similar concepts are represented by vectors that are close together in vector space.

Vector Databases (Section 4)

Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems. Vector databases store and search embeddings efficiently. They are commonly used in semantic search and RAG systems.

Chunking (Section 4)

Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance. Chunking divides large documents into smaller pieces before embedding. Proper chunking improves retrieval quality and context relevance.

Retrieval-Augmented Generation (Section 4)

RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context. RAG combines information retrieval with language generation. Relevant document chunks are retrieved and supplied to the language model as context.

Hybrid Search (Section 4)

Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance. Hybrid search combines vector similarity search with keyword-based methods such as BM25. This often improves retrieval performance.

Query Rewriting (Section 4)

Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems. Query rewriting transforms user questions into clearer standalone queries. This is especially useful for follow-up questions in conversational systems.

Reranking (Section 4)

Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking. Reranking evaluates retrieved documents and reorders them according to relevance. Cross-encoder models are frequently used for reranking.

Conversation Memory (Section 4)

Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations. Conversation memory allows systems to retain recent interactions and support multi-turn conversations.

Evaluation (Section 4)

RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction. RAG systems are evaluated using retrieval metrics, answer quality metrics, latency, and user satisfaction.

Future of RAG (Section 4)

Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities. Future RAG systems may include multimodal retrieval, agentic workflows, knowledge graphs, and advanced reasoning capabilities.